

# Statistics 452: Statistical Learning and Prediction

## Chapter 4, Part 2: Linear Discriminant Analysis Classifier

Brad McNeney

# Linear Discriminant Analysis (LDA) Classifier

# Logistic Regression vs LDA Classifieer

- ▶ Logistic regression models  $Pr(Y = 1|X = x) = p(x)$ .
  - ▶ More generally, polytomous regression models  $Pr(Y = k|X = x)$ .
- ▶ LDA models  $Pr(X = x|Y = k)$ , which relates to  $Pr(Y = k|X = x)$  via Bayes' rule:

$$Pr(Y = k|X = x) = \frac{Pr(X = x|Y = k)Pr(Y = k)}{\sum_j Pr(X = x|Y = j)Pr(Y = j)}.$$

# Notation and Terminology

- ▶ Notation:  $f_k(X) = Pr(X = x|Y = k)$ ,  $\pi_k = Pr(Y = k)$ ,  $p_k(x) = Pr(Y = k|X = x)$ , so that

$$p_k(x) = \frac{f_k(x)\pi_k}{\sum_j f_j(x)\pi_j}.$$

- ▶ Terminology:
  - ▶  $\pi_k$  is the *prior* probability of class  $k$
  - ▶  $p_k(x)$  is the *posterior* probability of class  $k$  given data  $x$ .

# Bayes Classifier

- ▶ We saw in Chapter 2 that the classifier with the lowest error rate is the one that chooses the class with the highest posterior probability.
- ▶ If we have good estimates,  $\hat{p}_k(x)$ , we can choose the class with the highest **estimated** posterior probability.
- ▶ Estimate  $p_k(x)$  by estimating
  - ▶  $\pi_k$ : If training sample is a population random sample, use the sample proportions.
  - ▶  $f_k(x)$ : LDA model is parametric, so we estimate  $f_k$  by estimating its parameters.

## LDA Classifier: Model for $f_k(x)$ when $p = 1$

- ▶ Assume  $f_k(x)$  is normal with mean  $\mu_k$  and variance  $\sigma_k^2$ .
  - ▶ Then estimating  $f_k$  amounts to estimating  $\mu_k$  and  $\sigma_k^2$ .
  - ▶ Further assume  $\sigma_k$  is the same for all  $k$  and let  $\sigma$  denote this common value.
- ▶ For a particular  $x$ , the class with highest posterior probability is the one with highest value of

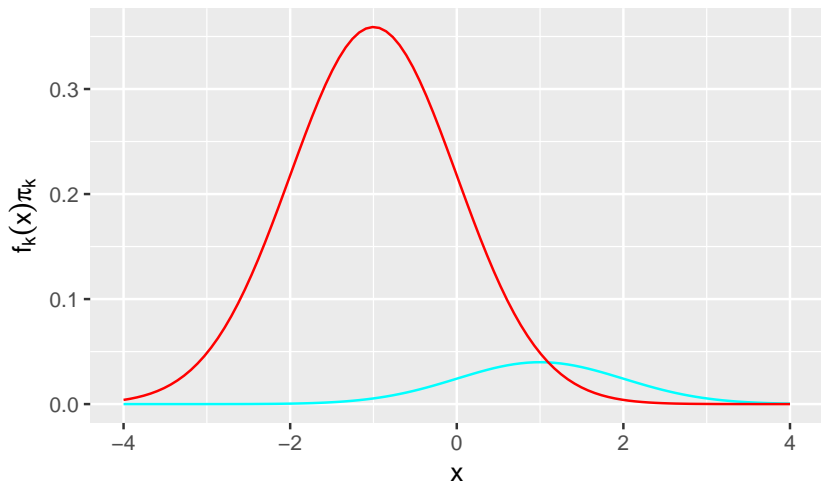
$$\delta_k(x) = x \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log \pi_k.$$

## Example

- Suppose two classes,  $\pi_1 = 0.1$ ,  $\pi_2 = 0.9$ ,  $\mu_1 = 1$ ,  $\mu_2 = -1$ ,  $\sigma = 1$ .

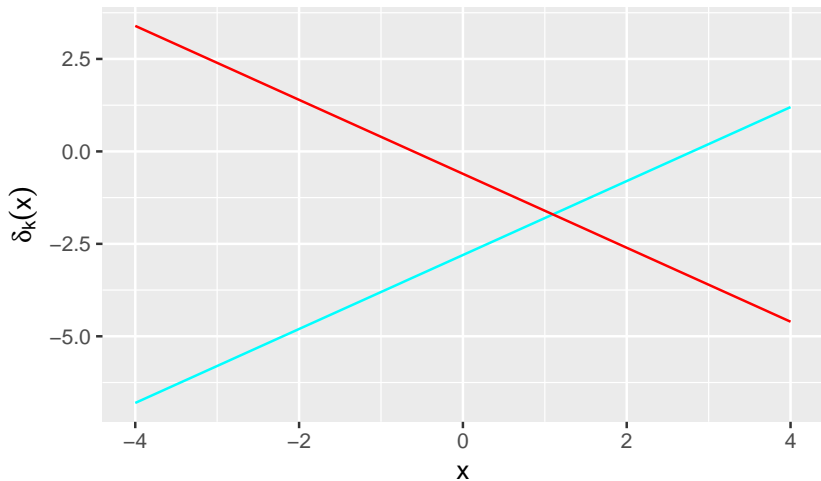
```
mu1<-1; mu2<-(-1); sigma<-1; pi1<-0.1;pi2<-0.9
x <-seq(from=-4,to=4,length=100)
f1 <- dnorm(x,mean=mu1,sd=sigma)
f2 <- dnorm(x,mean=mu2,sd=sigma)
delta <- function(x,mu,sigma,pi) {
  x*mu/sigma^2 - mu^2/(2*sigma^2) + log(pi)
}
delta1 <- delta(x,mu1,sigma,pi1)
delta2 <- delta(x,mu2,sigma,pi2)
dd <- data.frame(x=x,f1=f1,f2=f2,delta1=delta1,delta2=delta2)
```

```
library(tidyverse)
ggplot(dd,aes(x=x))+
  geom_line(aes(y=f1*pi1),color="cyan")+
  geom_line(aes(y=f2*pi2),color="red") +
  labs(y=expression(f[k](x)*pi[k]))
```





```
ggplot(dd,aes(x=x))+  
  geom_line(aes(y=delta1),color="cyan")+  
  geom_line(aes(y=delta2),color="red") +  
  labs(y=expression(delta[k](x)))
```



# Decision Boundary

- ▶ The point where the  $\delta_k$  lines cross is the decision boundary.
- ▶ Can show that the boundary is

$$\frac{\mu_1 + \mu_2}{2} + \frac{\sigma^2 \log(\pi_2/\pi_1)}{\mu_1 - \mu_2}.$$

- ▶ In this example the boundary is  
 $(1 + (-1))/2 + \log(.9/.1)/2 \approx 1.1$

```
(mu1 + mu2)/2 + sigma^2*log(pi2/pi1)/(mu1-mu2)
```

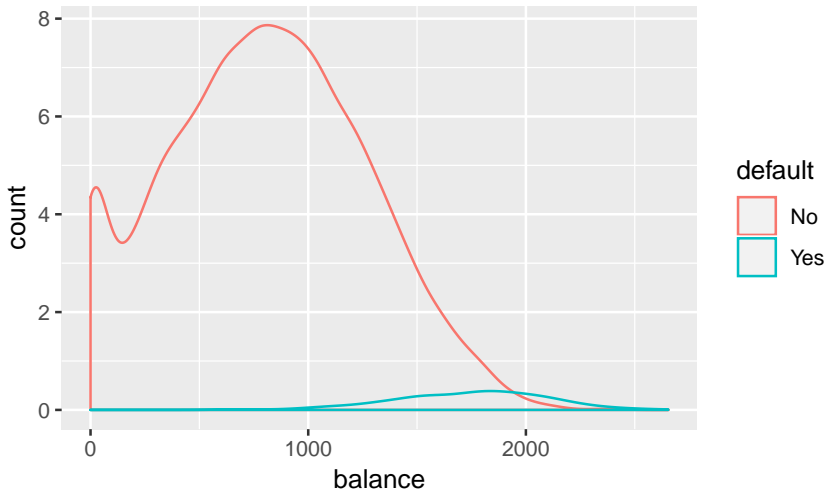
```
## [1] 1.098612
```

# Estimated Distributions

- ▶ Replace means and variance in  $\delta_k(x)$ 's with estimates (equation 4.15 of text) to get the *discriminant functions*  $\hat{\delta}_k$ .
- ▶ Illustrate with Default data.

```
library(ISLR)  
data(Default)
```

```
ggplot(Default,aes(x=balance,color=default)) +  
  geom_density(aes(y=..count..))
```



## Software Note

- ▶ Specifying `color=default` groups the data by default.
- ▶ The density geom plots the estimated densities  $\hat{f}_k(x)$  by default, but also computes the variable `..count..`, which for group  $k$  is  $\hat{f}_k(x)n_k$ .
- ▶ We are interested in  $\hat{f}_k(x)\hat{\pi} = \hat{f}_k(x)n_k/n$ , so `count` is proportional to what we want to plot.

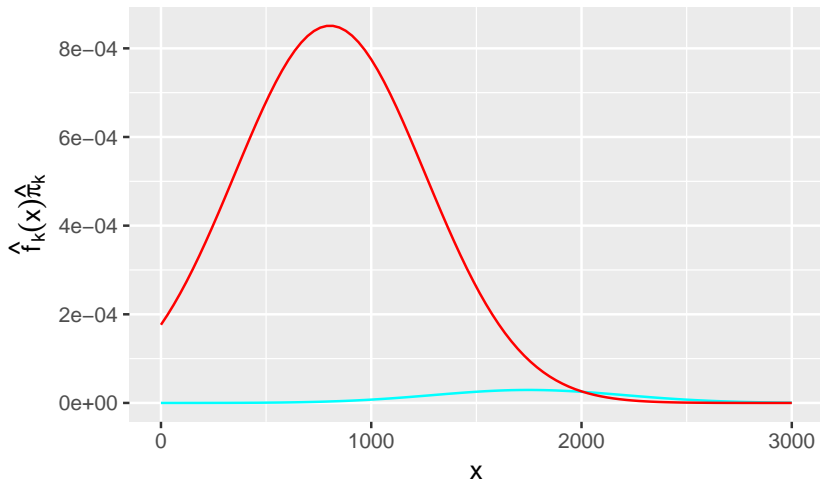
# Estimated Densities and Discriminant Functions

```
n <- nrow(Default); K <- 2
pi1 <- with(Default, mean(default=="Yes")); pi2 <- 1-pi1
defBalance <- with(Default, balance[default=="Yes"])
nodefBalance <- with(Default, balance[default=="No"])
mu1 <- mean(defBalance)
mu2 <- mean(nodefBalance)
sigma <- sqrt((sum((defBalance-mu1)^2) + sum((nodefBalance-mu2)^2))/(n-K))
with(Default, range(balance))
```

```
## [1]      0.000 2654.323
```

```
x <- seq(from=0, to=3000, length=100)
f1 <- dnorm(x, mean=mu1, sd=sigma)
f2 <- dnorm(x, mean=mu2, sd=sigma)
delta1 <- delta(x, mu1, sigma, pi1)
delta2 <- delta(x, mu2, sigma, pi2)
dd <- data.frame(x=x, f1=f1, f2=f2, delta1=delta1, delta2=delta2)
```

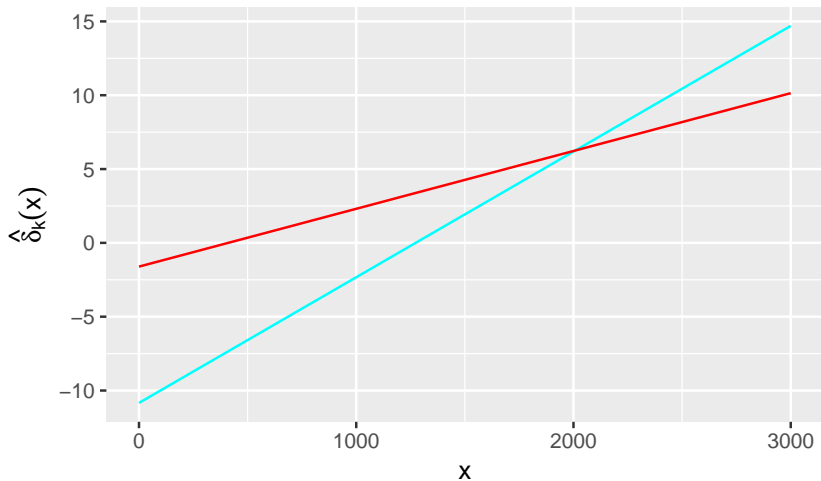
```
ggplot(dd,aes(x=x))+ geom_line(aes(y=f1*pi1),color="cyan")+  
  geom_line(aes(y=f2*pi2),color="red") +  
  labs(y=expression(hat(f)[k](x)*hat(pi)[k]))
```



```
min(x[f1*pi1>f2*pi2])
```

```
## [1] 2030.303
```

```
ggplot(dd,aes(x=x))+  
  geom_line(aes(y=delta1),color="cyan")+  
  geom_line(aes(y=delta2),color="red") +  
  labs(y=expression(hat(delta)[k](x)))
```





# Decision Boundary for Default Data

```
(mu1 + mu2)/2 + sigma^2*log(pi2/pi1)/(mu1-mu2)
```

```
## [1] 2008.584
```

# LDA Classifier in R

```
library(MASS)
ll <- lda(default ~ balance, data=Default)
preds <- predict(ll)
head(preds$posterior)
```

```
##           No           Yes
## 1 0.9972130 0.002786981
## 2 0.9958358 0.004164240
## 3 0.9865931 0.013406929
## 4 0.9988882 0.001111757
## 5 0.9963955 0.003604464
## 6 0.9933487 0.006651334
```

```
head(preds$class)
```

```
## [1] No No No No No No
## Levels: No Yes
```

## LDA Classifier for $p > 1$

- ▶ Same basic idea of classifying based on highest posterior probabilities.
- ▶ Generalize the model for  $f_k(X)$  from a Gaussian distribution to a *multivariate* Gaussian distribution.
- ▶ The mean in class  $k$  is now a vector  $\mu_k$  and the variance is a matrix  $\Sigma$ .
- ▶ Details on the multivariate density and the resulting  $\delta_k(x)$  are given on page 143 of the text.

# LDA Classifier Using all Default Data

```
ll <- lda(default ~ student + balance + income, data=Default)
preds <- predict(ll)
head(preds$posterior)
```

##		No	Yes
## 1	0.9967765	0.003223517	
## 2	0.9973105	0.002689531	
## 3	0.9852914	0.014708600	
## 4	0.9988157	0.001184329	
## 5	0.9959768	0.004023242	
## 6	0.9957918	0.004208244	

# Predictions vs True Default Status

- ▶ Tabulating predictions and true default status gives the following “confusion matrix”:

```
Default <- data.frame(Default, predicted = preds$class)
xtabs( ~ predicted + default, data=Default)
```

```
##           default
## predicted   No   Yes
##           No  9645  254
##           Yes   22   79
```

- ▶ (Slight difference with results in text.)
- ▶ 9724 correctly predicted, 276 incorrectly predicted, so error rate is 2.76%.
- ▶ But 254/333, or 76% of defaulters missed.

# Classification Error Rates

- ▶ Terminology is from diagnostic testing in medicine.
  - ▶ E.G., pap smear to screen for cervical cancer
- ▶ The test can be positive ( $T^+$ ) or negative ( $T^-$ ), and an individual may have the disease ( $D^+$ ) or not ( $D^-$ ):

		Disease Status	
		$D^-$	$D^+$
Test Result	$T^-$	True Negative (TN)	False Negative (FN)
	$T^+$	False Positive (FP)	True Positive (TP)
		TN+FP	FN+TP

- ▶ Replace test with prediction and true disease classification with true classification.

# Sensitivity, Specificity, True- and False-Positive Rates

- ▶ The observed true positive rate (TPR) is the proportion of  $D^+$  who are TP, or  $TP/(TP+FN)$ .
  - ▶  $TP/(TP+FN)$  is an estimate of  $P(T^+ | D^+)$ , known as the *sensitivity* of the test.
  - ▶ Sensitivity is only 79/333 or 23.7% for Default data.
- ▶ The observed true negative rate is the proportion of  $D^-$  who are TN, or  $TN/(TN+FP)$ .
  - ▶  $TN/(TN+FP)$  is an estimate of the true negative rate  $P(T^- | D^-)$ , known as the *specificity* of the test. Specificity is 9645/9667 or 99.8% for Default data.
  - ▶ The complement,  $FP/(TN+FP)$ , is an estimate of the false positive rate (FPR)  $P(T^+ | D^-)$ . FPR is 22/9667 or 0.2% for Default data.

# Receiver Operating Characteristic (ROC) Curves

- ▶ The classification of a subject with  $X = x_0$  as  $T^+$  or  $T^-$  is made by comparing the posterior probability  $Pr(T^+|X = x_0)$  to a threshold  $t = 1/2$ ; e.g.,  $> t$  classify as  $T^+$ ,  $\leq t$  classify as  $T^-$ .
- ▶ Lowering the threshold will increase the number of  $T^+$ , and hence the TPR,  $P(T^+ | D^+)$ , and the FPR,  $P(T^+ | D^-)$ .
- ▶ The ROC curve is a plot plot of the TPR *versus* FPR as we vary  $t$ .
  - ▶ Try to get a sense of the trade-off between true- and false-positive rates, and possibly choose an “optimal”  $t$ .

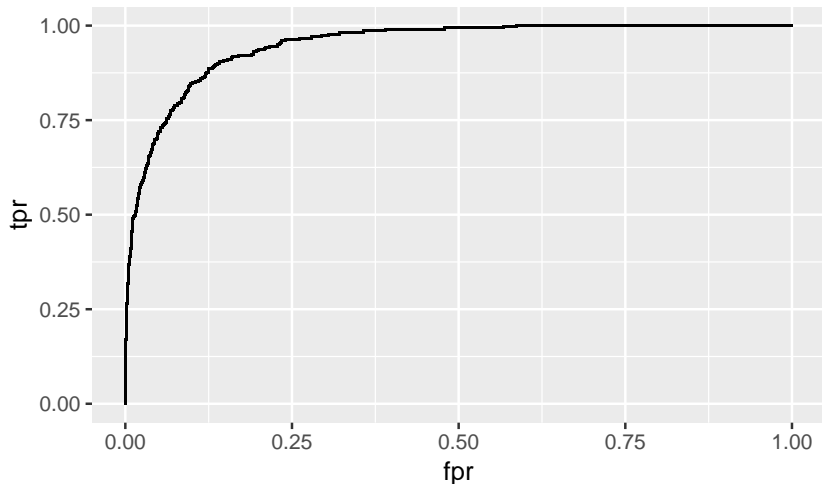


# ROC for LDA Classifier on Default Example

```
library(AUC) # for the roc() function; install.packages("AUC") to install
library(broom) # for the tidy() function
posteriorYes <- preds$posterior[, "Yes"]
trueYes <- # require a binary factor with 1=Yes
  (Default$default=="Yes") %>% as.numeric() %>% factor()
ROCRes <- roc(posteriorYes, trueYes)
tidyROCRes <- tidy(ROCRes)
```

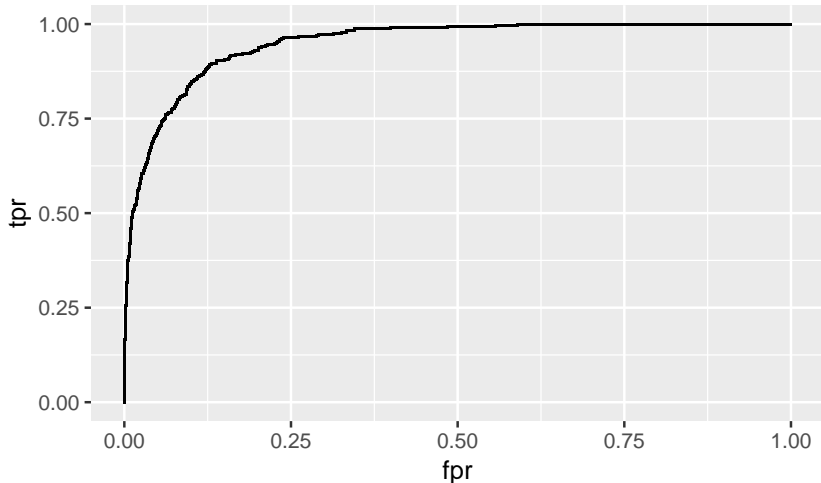
# ROC Curve

```
ggplot(tidyROCres, aes(x=fpr, y=tpr)) + geom_point(pch=".")
```



# ROC Curve for Logistic Regression

```
lfit <- glm(default ~ ., data=Default, family=binomial())  
posteriorYesLogistic <- fitted(lfit)  
ROClogist <- tidy(roc(posteriorYesLogistic, trueYes))  
ggplot(ROClogist, aes(x=fpr, y=tpr)) + geom_point(pch=".")
```



# ROC Interpretation

- ▶ The points on the plot represent the FPR/TPR combinations for each threshold.
- ▶ The ideal test threshold would yield a TPR of one and a FPR of zero, and so would appear in the upper-left corner of the ROC plot.
  - ▶ We usually select the threshold that is closest to the upper-left corner.
  - ▶ No obvious “best” threshold in this example. Perhaps the threshold that gives TPR of about 0.875 and FPR of about 0.125.

```
tidyROCres %>%  
  filter(fpr >= 0.1249, fpr <= 0.1251)
```

```
## # A tibble: 3 x 3  
##   cutoffs   fpr   tpr  
##   <dbl> <dbl> <dbl>  
## 1  0.0457 0.125 0.877  
## 2  0.0457 0.125 0.877  
## 3  0.0457 0.125 0.880
```

# Re-classify

```
n <- nrow(Default); thresh <- 0.0457
dclass <- rep("No",n); dclass[posteriorYes>thresh] <- "Yes"
Default <- data.frame(Default,prednew = dclass)
xtabs(~ prednew + default, data=Default)
```

```
##           default
## prednew    No  Yes
##      No  8458   41
##      Yes 1209  292
```

- ▶ Sensitivity is much better (about 87.7%).
- ▶ Many more false-positives (1209)

# Area Under Curve (AUC)

- ▶ The area under the ROC curve is a measure of the overall performance of the classifier.
- ▶ If TPR jumps to near one with FPR remaining low, the performance is good.
  - ▶ This would give AUC near 1.
- ▶ A poor classifier would do no better than the null classifier, which for threshold  $t$  randomly assigns  $t \times 100\%$  to be the success category and has  $AUC=1/2$ .

```
auc(ROCres)
```

```
## [1] 0.9495202
```

# Quadratic Discriminant Analysis (QDA)

- ▶ Back to the  $p = 1$  case.
- ▶ If we let the variances differ by class  $k$ , then it can be shown that the Bayes classifier assigns an observation with  $X = x$  to the class with the largest

$$\delta_k(x) = -\frac{1}{2} \frac{(x - \mu_k)^2}{\sigma^2} - \frac{1}{2} \log \sigma_k^2 + \log \pi_k,$$

which is a quadratic function of  $x$ .

- ▶ See the text, page 149, for the multivariate version.

# QDA for the Default Data

```
qq <- qda(default ~ student + income + balance, data=Default)
preds <- predict(qq)
Default <- data.frame(Default, predicted = preds$class)
xtabs( ~ predicted + default, data=Default)
```

```
##           default
## predicted   No   Yes
##           No  9645  254
##           Yes   22   79
```

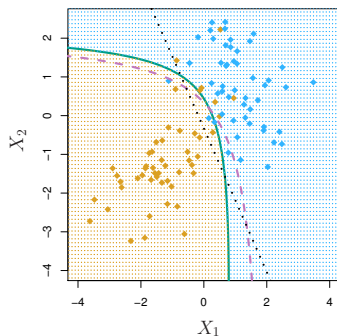
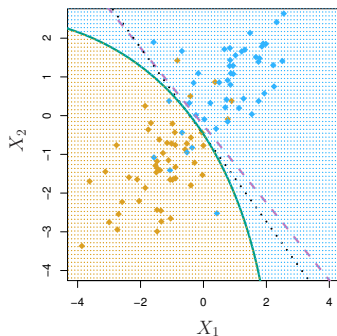
- Same as LDA for these data.



## QDA vs LDA: Bias vs Variance

- ▶ When should we assume the  $K$  classes have a common variance matrix?
- ▶ LDA has fewer parameters and so lower variance.
  - ▶ When  $p$  is large relative to  $n$ , the number of parameters in QDA can become large, and variance of the posterior probabilities large.
- ▶ On the other hand, if the variances really **are** different, the LDA posterior probabilities can be biased.
- ▶ Recommendation from the text: If the training data set is large, use QDA. If not, use LDA.

# QDA vs LDA: Decision Boundaries



- Text, Figure 4.9:  $p = 2$ , Bayes (purple dashed), LDA (black dotted), QDA (green solid).

# Comparison of Classification Methods

- ▶ Performance depends on the nature of the true decision boundary (curve that separates the two classes).
- ▶ One can show that logistic regression leads to linear decision boundaries, just like LDA
  - ▶ Difference is in how the methods estimate the parameters of the boundaries.
  - ▶ Often give very similar results, as in the Default data.
- ▶ QDA gives quadratic boundaries.
- ▶ The KNN classifier is non-parametric and so there are no restrictions on the decision boundary.
- ▶ Performance depends on the nature of the true decision boundary.