

Statistics 452: Statistical Learning and Prediction

Chapter 7, Part 3: Smoothing Splines, Local Regression

Brad McNeney

Smoothing Splines

Smoothing Splines Overview


- ▶ Smoothing splines are an alternative approach to devising a smooth curve.
- ▶ Though the derivation is different, we end up with a natural cubic spline, with knots *at each observed data point*.
 - ▶ The estimated coefficients of the basis functions using least squares would interpolate the data points.
 - ▶ To avoid over-fitting, a shrinkage penalty is added to the RSS. The penalty is comprised of a tuning parameter times a penalty term.

Smoothness

- ▶ We seek a *smooth* function $g(x)$ that fits the data well.
 - ▶ Want $RSS = \sum_{i=1}^n (y_i - g(x_i))^2$ small.
- ▶ By smooth we mean not too “rough” (wiggly)
 - ▶ If RSS is the objective function, the curve g will interpolate the data and will be very rough.
 - ▶ Illustrate by fitting a natural cubic spline to simulated data on “age” and “wage”

```
set.seed(1)
age <- scale(18:80); betas <- c(2,-2,2,-2)
f <- poly(age,degree=4)%*% betas
wage <- f + rnorm(length(age))
sfrit0 <- smooth.spline(age,wage,lambda=0)
newage<-seq(from=min(age),to=max(age),length=1000)
pwage0 <- predict(sfrit0,newage)
```

```
plot(age,wage)
lines(pwage0$x,pwage0$y,col="blue") # fitted curve
lines(age,f,col="red") # true curve
```



Penalized Objective Function

- ▶ The criterion function to minimize is


$$\sum_{i=1} (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt. \quad (1)$$

- ▶ Like ridge regression or the lasso, the objective function is of the form RSS + penalty, and we have a tuning, or smoothing parameter λ , but the penalty function is now $\int g''(t)dt$.
 - ▶ $g''(t)$ is the curvature of g at t .
 - ▶ $\int g''(t)^2 dt$ is a measure of the total curvature.
- ▶ It can be shown that the minimizer of (1) is a natural cubic spline with knots at the unique observed data points, with the coefficients of the basis functions *shrunk* towards zero.
 - ▶ The degree of shrinkage is controlled by λ .

```

sfit1 <- smooth.spline(age,wage,spar=1) #spar= c1+c2*log(lambda)
pwage1 <- predict(sfit1,newage)
plot(age,wage)
lines(pwage0$x,pwage0$y,col="blue") # fitted curve with lambda=0
lines(pwage1$x,pwage1$y,col="green")
lines(age,f,col="red") # true curve

```



Choosing the Smoothing Parameter

- ▶ Could use cross validation to select λ .
 - ▶ How many folds?
- ▶ It turns out there is a very simple formula for the CV estimate of test error when using leave-out-one CV (LOOCV):

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{g}_{\lambda}(x_i)^{(-i)})^2 = \frac{1}{n} \sum_{i=1}^n \left(\frac{y_i - \hat{g}_{\lambda}(x_i)}{1 - \{S_{\lambda}\}_{ii}} \right)^2 \quad (2)$$

where

- ▶ $\hat{g}_{\lambda}(x_i)^{(-i)}$ is the fitted value for x_i when (x_i, y_i) is held out of the training set,
- ▶ $\hat{g}_{\lambda}(x_i)$ is the fitted value for x_i when all of the data are used to fit g ,
- ▶ $\{S_{\lambda}\}_{ii}$ is the i th diagonal entry of the “smoother matrix”.

The Smoother Matrix

- ▶ The smoother matrix S_λ is the matrix that turns the response y into the smooth \hat{g}_λ ; i.e., $\hat{g}_\lambda(x) = S_\lambda y$, where x is the vector of explanatory variable values and $\hat{g}_\lambda(x)$ is the vector of $\hat{g}_\lambda(x_i)$ values.
- ▶ In linear regression, the smoother matrix is called the hat matrix and is denoted H .
 - ▶ H turns the response y into \hat{y} ; i.e., $\hat{y} = Hy$,
 - ▶ The sum of the diagonal elements of H is the number of regression coefficients $p + 1$.
 - ▶ The diagonal elements are the hat values, or leverages h_i .
- ▶ Thus the $CV_{(n)}$ estimate for the smoothing spline is analogous to least squares, with the diagonal elements of H replaced by those of S_λ .


Effective Degrees of Freedom

- ▶ In linear regression, the model df $p + 1$ can be shown to equal the sum of the diagonal elements of H .
- ▶ By analogy, the sum of the diagonal elements of the smoother matrix is referred to as the *effective degrees of freedom* and is denoted df_λ .
 - ▶ One can show that as λ ranges from 0 to ∞ , df_λ ranges from n to 2.
 - ▶ Can treat df_λ as the smoothing parameter and select its value by CV.

```
sfitCV <- smooth.spline(age,wage,cv=TRUE)
sfitCV$df
```

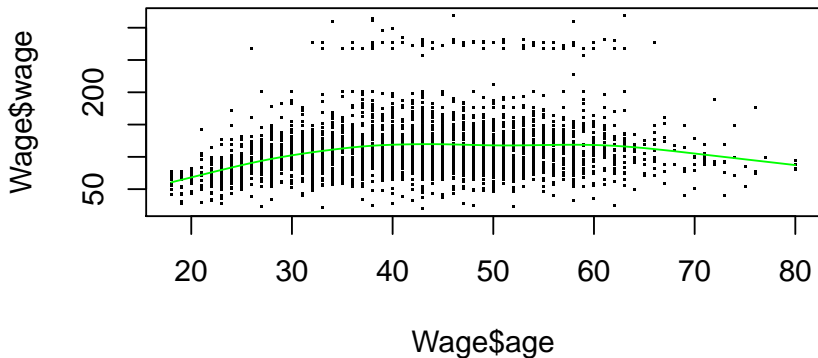
```
## [1] 5.769452
```

```
pwageCV <- predict(sfitCV,newage)
plot(age,wage)
lines(pwageCV$x,pwageCV$y,col="green")
lines(age,f,col="red") # true curve
```



Example

```
library(ISLR); data(Wage)
plot(Wage$age,Wage$wage,pch=".")
sfit <- smooth.spline(Wage$age,Wage$wage,cv=TRUE)
pwage <- predict(sfit)
lines(sfit,col="green") # can plot output of smooth.spline directly
```



Local Regression

Local Regression

- ▶ A variation on KNN: instead of fitting a constant over a neighborhood, fit a weighted regression.
 - ▶ weighted means points in the neighborhood closest to the point of prediction are up-weighted.
 - ▶ the regression could be constant, linear or quadratic.
- ▶ The neighborhood size is referred to as the “span” s .

Local Linear Regression Algorithm

- ▶ Select a span s and a weight function $K(x_i, x_0)$.
- ▶ For $X = x_0$:
 1. Extract the nearest $s * n$ points to x_0 to train the local regression.
 2. Assign weight $K_{0i} = K(x_i, x_0)$ to each neighbor.
 3. Fit a weighted least-squares regression; that is, find the $\hat{\beta}_0$ and $\hat{\beta}_1$ that minimize

$$\sum_{i=1}^n K(x_i, x_0)(y_i - [\beta_0 + \beta_1 x_i])^2$$

4. $\hat{f}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0$.


```

plotfit <- function(fit,dat,newdat){
  preds <- data.frame(newdat,
    predict(fit,newdata=newdat,se=TRUE))
  preds <- mutate(preds, # approx. CIs from the SEs
    lwr = preds$fit-2*preds$se,upr = preds$fit+2*preds$se)
  ggplot(dat,aes(x=age,y=wage)) + geom_point(alpha=0.1) +
    geom_ribbon(aes(x=age,y=fit,ymin=lwr,ymax=upr),
      data=preds,fill="blue",alpha=.2) +
    geom_line(aes(y=fit),data=preds,color="blue")
}


```



```
library(ggplot2); library(dplyr)
newdat <- data.frame(age=seq(from=min(Wage$age),to=max(Wage$age),length=100))
fit2 <- loess(wage~age,span=0.2,data=Wage)
plotfit(fit2,Wage,newdat)
```



```
fit5 <- loess(wage~age,span=0.5,data=Wage)  
plotfit(fit5,Wage,newdat)
```



Extensions to Higher Dimension

- ▶ For $p > 1$ explanatory variables, we can generalize local regression.
 - ▶ Choose neighborhoods based on X_1, \dots, X_p .
 - ▶ Fit a multiple linear regression.
- ▶ However, it has been observed that local regression performs poorly for $p > 3$ because there are few points close to x_0 (recall homework 2).