

Statistics 452: Statistical Learning and Prediction

Chapter 9: Support Vector Machines

Brad McNeney

Overview


- ▶ The context is a binary classification problem.
 - ▶ We will discuss two classification rules that incrementally lead to the support vector machine.
1. Maximal margin classifier: If there exists a linear boundary in the space of explanatory variables that separates the classes in the training data, use this boundary as a classifier.
 - ▶ If one boundary exists, there will be many. Choose one that maximizes the “margin”, a minimum distance between predictors and the boundary.
 - ▶ Simple classifier, but requires that a linear boundary exist.
 2. Support vector classifier: Weaken the requirement of complete separation to a linear boundary that **best** separates the classes.
 - ▶ Better, but a linear boundary may not be the best choice.
- ▶ Support vector machine: Weaken the requirement of a linear boundary, and add a trick to prevent computations from becoming prohibitive.

Linear Boundaries

- ▶ Linear boundaries can be represented as a “hyperplane”, which is a $p - 1$ -dimensional subspace of p dimensions specified by a constraint

$$f(X) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p = 0.$$



- ▶ Classify as blue class if $f(X) > 0$ and red if $f(X) < 0$.



(Text, Figure 9.1)

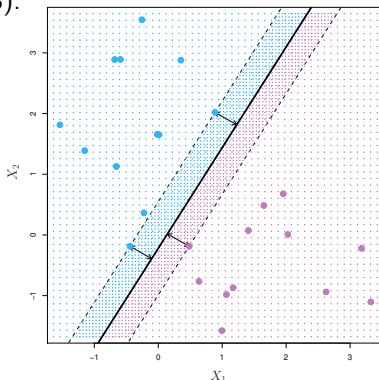
Possible Hyperplanes

- ▶ When a separating hyperplane exists, there are many.
- ▶ Text, Figure 9.2:
 - ▶ Left panel, three separating hyperplanes
 - ▶ Right panel, one hyperplane and resulting decision rule



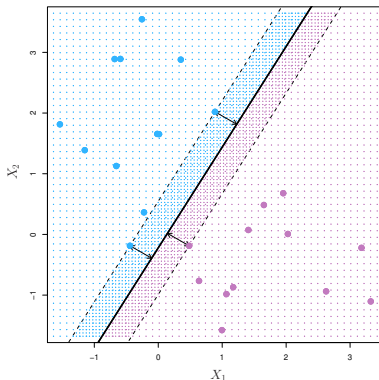
Maximal Margin Classifier

- ▶ The maximal margin hyperplane is the separating hyperplane farthest from the training observations.
 - ▶ For each possible hyperplane, find the margin: The minimum of perpendicular distances between each point and the hyperplane.
 - ▶ Choose the hyperplane with the largest margin as the boundary (Fig. 9.3).



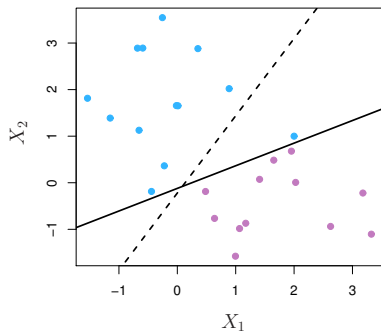
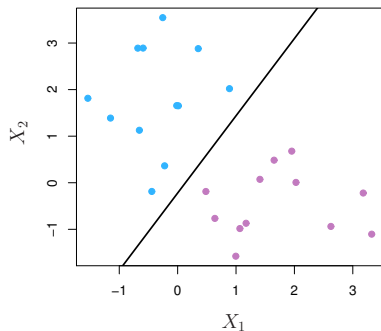
Support Vectors

- ▶ The support vectors are the points (p -dimensional vectors) that determine the maximal margin hyperplane.
 - ▶ Support in the sense that if they move, so does the hyperplane.
 - ▶ Support vectors in the diagram are two blue points and one red with perpendicular distances indicated on the figure.
 - ▶ The classifier depends on the data through the support vectors **only**.



Sensitivity to Support Vectors

- ▶ Classifier depends on data through the support vectors **only**.
- ▶ Adding/deleting a support vector can drastically change the maximal margin hyperplane.
- ▶ Example maximal margin hyperplane before (left panel) and after (right panel) adding a new support vector (text, Fig 9.5):

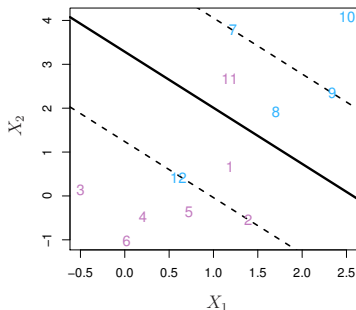
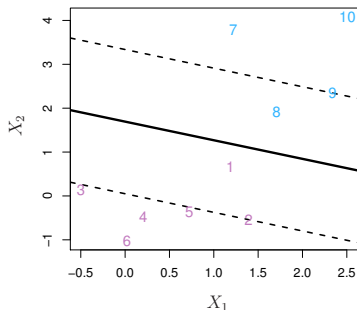


A Budget of Margin Violations

- ▶ Notice that the margin in right panel of the previous example is very small.
- ▶ We might instead seek a larger margin, but allow some points to be on the wrong side.
 - ▶ Allow vector i to be a distance ϵ_i on the wrong side of the margin, as long as $\sum_{i=1}^n \epsilon_i \leq C$, for some budget C .
 - ▶ By allowing points on the wrong side of the margin, and even wrong side of the hyperplane, we can accommodate the case where the vectors are not separable by a hyperplane.
- ▶ This is the support vector classifier.
 - ▶ The budget C is a tuning parameter.
- ▶ The support vectors are those on the margin or on the wrong side of the margin.
 - ▶ One can show that only these support vectors affect the choice of hyperplane and hence classification rule.

Illustration of Support Vector Classifier

- ▶ In the left panel of the following Figure (Fig 9.6), one observation from each class is on the wrong side of the margin, but not the hyperplane.
- ▶ In the right panel, two more points are added that are also on the wrong side of the hyperplane.
- ▶ The support vectors are those on the margin or on the wrong side of the margin.



Choosing the budget C

- ▶ Bias-variance tradeoff:
 - ▶ Small C means we require narrow margins, and are potentially over-fitting the data. Should have low bias but high variance
 - ▶ Large C means we allow wide margins and are not fitting the data as aggressively. Should have higher bias, but lower variance.
- ▶ Use cross-validation to select C .

Example Support Vector Classifier for the Heart Data

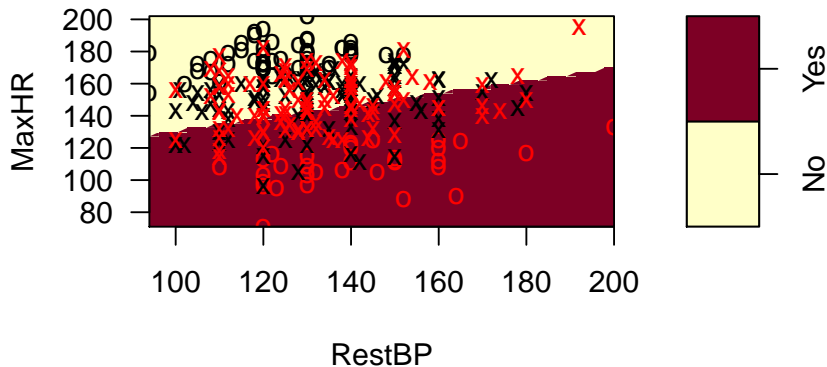
- ▶ The `svm()` function from the `e1071` package fits support vector machines.
 - ▶ Set argument `kernel=linear` for the support vector classifier.

```
uu <- url("http://faculty.marshall.usc.edu/gareth-james/ISL/Heart.csv")
Heart <- read.csv(uu,row.names=1)
Heart <- na.omit(Heart)
library(dplyr)
Heart <- select(Heart,AHD,MaxHR,RestBP)
library(e1071)
svc.heart <- svm(AHD ~ .,type="C-classification",cost=1,
                 data=Heart,kernel="linear")
```

Visualize the Classifier

```
plot(svc.heart,Heart)
```

SVM classification plot



- ▶ (Zoom for better view) Classes are colour-coded; O means on right side and X means wrong side of margin.

Choose the cost by CV

- ▶ The `tune()` function does the CV over a user-supplied grid of costs.

```
set.seed(123)
tune.heart <- tune(svm,AHD ~ .,data=Heart,kernel="linear",
                  ranges=list(cost=c(10^{-3:3})))
summary(tune.heart)$performances
```

```
##      cost      error dispersion
## 1 1e-03 0.4610345 0.08855707
## 2 1e-02 0.3265517 0.06894093
## 3 1e-01 0.3095402 0.06761133
## 4 1e+00 0.3029885 0.07597847
## 5 1e+01 0.3064368 0.07361748
## 6 1e+02 0.3064368 0.07361748
## 7 1e+03 0.3031034 0.07949254
```

- ▶ According to these CV results, we should use cost 1.

Non-linear Decision Boundaries

- ▶ By expanding the features to include polynomial terms, the support vector classifier will have non-linear decision boundaries in the original feature space.
 - ▶ In the expanded feature space, say $X_1, X_1^2, \dots, X_p, X_p^2$, the boundary will be a curve of the form

$$f(X) = \beta_0 + \beta_1 X_1 + \beta_2 X_1^2 + \dots + \beta_{2p-1} X_p + \beta_p X_p^2 = 0.$$

- ▶ We could consider higher-order polynomials, but eventually we would have so many features that computation time would become a problem.
- ▶ The support vector machine uses a computational trick to allow non-linear decision boundaries without prohibitive computation.
 - ▶ To describe the trick, recast the computation for the support vector classifier in terms of “kernels”.

Support Vector Classifier *via* Kernels

- ▶ It turns out (e.g., ESL, Section 12.3) that the classification rule for a point x depends on the sign of

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i K(x, x_i)$$

where

- ▶ $K(x, x_i) = \sum_{j=1}^p x_j x_{ij}$ and
- ▶ β_0 and the α_i 's depend on the data only through the $n(n-1)/2$ values of $K(x_i, x_{i'})$.
- ▶ The function $K(\cdot, \cdot)$ is called the linear kernel.

Support Vector Machine

- ▶ We can extend the approach by choosing other kernel functions $K(\cdot, \cdot)$.
 - ▶ In general, kernels are functions that measure **similarity** between two feature vectors (high value for similar vectors, low value for dissimilar vectors).
- ▶ Examples of alternative kernels include
 - ▶ the polynomial kernel, for given degree d :

$$K(x_i, x_{i'}) = \left(1 + \sum_{j=1}^p x_{ij} x_{i'j}\right)^d$$

- ▶ and the radial kernel for given $\gamma > 0$:

$$K(x_i, x_{i'}) = \exp\left(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2\right).$$

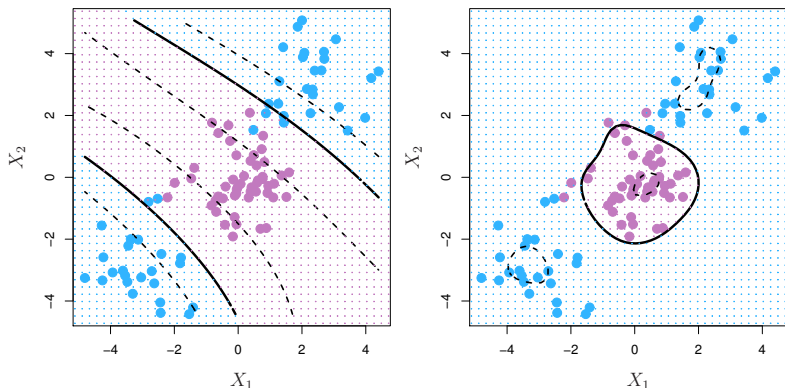
- ▶ The extension of the support vector classifier to non-linear boundaries, with an expanded feature space and computations done *via* kernels is called the support vector machine.

Computation for the SVM

- ▶ Key point: No matter how large the dimension of the expanded feature space, the computations only rely on the data through $n(n - 1)/2$ values of $K(x_i, x_{i'})$.

Example Decision Boundaries for Polynomial and Radial Kernels

- Figure 9.9. Left panel is from the polynomial kernel with $d = 3$; right panel is from a radial kernel (value of γ not given).

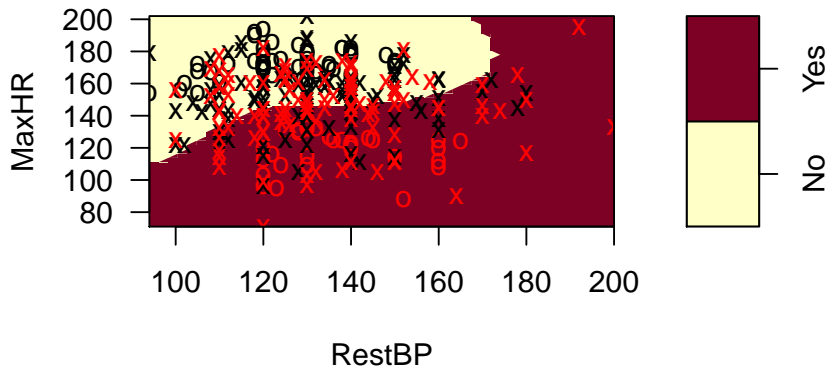


SVM for Heart Data

- Use a radial kernel.

```
svm.heart <- svm(AHD ~ ., type="C-classification", cost=1,  
                 data=Heart, kernel="radial", gamma=1/2)  
plot(svm.heart, Heart)
```

SVM classification plot



Choose the Cost and γ by CV

- ▶ The `tune()` function can do the CV over a grid of costs **and** other tuning parameters such as γ for the radial kernel.

```
set.seed(123)
tune2.heart <- tune(svm,AHD ~ .,data=Heart,kernel="radial",
                   ranges=list(cost=c(10^{-3:3}),
                               gamma=c(.5,1:4)))
perf<-summary(tune2.heart)$performances
perf[which.min(perf$error),]
```

```
##      cost gamma      error dispersion
## 10  0.1      1 0.2929885 0.08907935
```

Classifications

- ▶ Here we predict the training data.
 - ▶ In your weekly exercises you will split the Heart data into training and test sets and will predict the test set.

```
preds <- predict(svm.heart)
table(preds,Heart$AHD)
```

```
##
## preds  No  Yes
##   No  127  47
##   Yes   33  90
```

- ▶ About a 27% misclassification.
 - ▶ Looks pretty bad, but we're using just MaxHR and RestBP at this point.

SVMs with More than Two Classes

- ▶ No natural extension of the separating hyperplanes idea to multiple classes.
- ▶ Instead, use one-versus-one or one-versus-all classification.

One-versus-One Classification

- ▶ If there are K classes, there are $K(K - 1)/2$ pairwise comparisons.
- ▶ Can train $K(K - 1)/2$ SVMs.
- ▶ For a test observation, get $K(K - 1)/2$ predictions.
 - ▶ Each prediction is like a game in a tournament.
 - ▶ The final classification for the test observation is the winner of the tournament; i.e., the class that the test observation was most frequently assigned to.

One-versus-All Classification

- ▶ For each class, train a SVM to classify that class versus all others pooled together.
 - ▶ End up with K classifiers.
- ▶ The classification of a test observation is class with highest confidence; i.e., on the right side of the decision boundary, and farthest from the boundary.

Example: Gene Expression Data

- ▶ Four tumor types ($K = 4$) measured on 63 training and 20 test tumors.
- ▶ For each tumor, there are 2308 measurements of “gene expression”.
 - ▶ With this many features relative to the number of observations, non-linear kernels may provide **too much** flexibility – use linear.
- ▶ Classify tumor type based on expression measurements.
- ▶ `svm()` uses one-vs-one classification

```
library(ISLR)
data(Khan)
# help(Khan)
dat <- data.frame(y=as.factor(Khan$ytrain),
                  Khan$xtrain)
head(names(dat))
```

```
## [1] "y"  "X1" "X2" "X3" "X4" "X5"
```

```
library(e1071)
fit <- svm(y~., data=dat, kernel="linear", cost=10)
pp <- predict(fit, newdata=data.frame(Khan$xtest))
table(pp, Khan$ytest)
```

```
##
## pp  1 2 3 4
##    1 3 0 0 0
##    2 0 6 2 0
##    3 0 0 4 0
##    4 0 0 0 5
```

2/20

```
## [1] 0.1
```

The SVM and Logistic Regression

Loss + Penalty Formulation of the Support Vector Classifier

- ▶ Let \mathbf{X} and \mathbf{y} denote the matrix of X 's and vector of y 's, respectively and

$$f(\mathbf{X}; \beta) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$$

be the decision boundary.

- ▶ One can show that the criterion function that the support vector classifier minimizes to estimate f is of the form

$$\min_{\beta} \{L(\mathbf{X}, \mathbf{y}, \beta) + \lambda P(\beta)\}$$

where $L()$ is the so-called hinge loss function

$$L(\mathbf{X}, \mathbf{y}, \beta) = \sum_{i=1}^n \max[0, 1 - y_i f(x_i; \beta)],$$

$P(\beta) = \sum_{j=1}^p \beta_j^2$ is the ℓ_2 penalty function and λ is a tuning parameter.

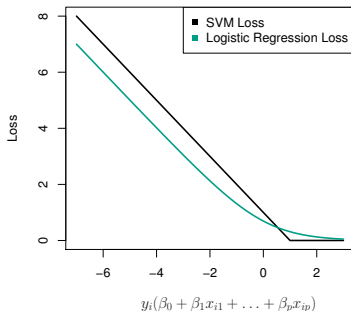
Hinge *versus* Logistic Loss

- ▶ The hinge loss function is similar to the logistic loss function

$$\sum_{i=1}^n \log(1 + e^{-y_i f(x_i; \beta)}),$$

used in logistic regression.

- ▶ Recall that logistic regression is fit by maximum likelihood.
- ▶ ML amounts to minimizing a negative-log-likelihood loss.
- ▶ Loss, as written above, is for outcomes coded as $-1/1$.



Support Vector Classifier/Machine *versus* Logistic Regression

- ▶ Conclude that SV Classifier is similar to logistic regression penalized with an ℓ_2 penalty.
- ▶ Can further argue that the SV Machine is similar to ℓ_2 -penalized logistic regression with non-linear functions of the predictors.